

GDSQ-VLA: GEOMETRY- AND DISTRIBUTION-SENSITIVE LAYER SELECTION FOR POST-TRAINING QUANTIZATION OF VISION-LANGUAGE-ACTION MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

Post-training quantization can reduce the static storage of vision-language-action (VLA) policies, but uniform precision overlooks the different numerical sensitivity of language reasoning and iterative action generation. We present GDSQ-VLA, a training-free mixed-precision framework that combines geometry and distribution preservation, paired action sensitivity, stability guards, and complete-configuration adjudication to allocate W4/FP16 layers under a frozen byte budget. The final method additionally uses a calibration-gated runtime correction selector: from model-level calibration statistics alone, it chooses the uncorrected policy, attention-temperature modulation (ATM), or output-head balancing (OHB), never their uncalibrated combination. The selector does not use task identity, rollout outcomes, or runtime success feedback. On RoboCasa365, 256-observation equivalence tests are bitwise exact for both evaluated policy families, allowing the GR00T selector’s baseline decision to reuse its frozen 50-task evaluation. This final GR00T configuration obtains 50.8% task-macro success at a theoretical $1.99\times$ candidate-component compression, compared with 55.1% for FP16. The audited $\pi_{0.5}$ selector obtains 27.5% task-macro SR over all 50 tasks and improves over its uncorrected static mask by 2.2 points (95% task CI [0.6, 3.9]). The current evidence supports the mechanism and a favorable GR00T storage–accuracy point, but not yet same-budget superiority across architectures.

1 INTRODUCTION

Vision-language-action (VLA) models connect visual perception and language-conditioned reasoning to continuous robot control, allowing one policy to execute diverse instructions across scenes and embodiments Brohan et al. (2023b); Driess et al. (2023); Brohan et al. (2023a); Octo Model Team et al. (2024); Kim et al. (2024). Diffusion- and flow-based action heads improve continuous action generation Chi et al. (2023); Liu et al. (2025a); Black et al. (2024); Physical Intelligence et al. (2025); NVIDIA et al. (2025), but large multimodal backbones and iterative denoising make deployment storage and memory bandwidth important constraints. Post-training quantization (PTQ) is appealing because it changes numerical precision without retraining the policy or redesigning its operator graph Xiao et al. (2023); Lin et al. (2024a).

VLA quantization is difficult for two related reasons. First, numerical sensitivity is heterogeneous: an error in a language layer changes multimodal conditioning, whereas an error in an action layer can accumulate across denoising steps. Uniform low precision ignores this structure, while a fixed architecture-level protection rule cannot adapt to a new checkpoint. Second, a mixed-precision mask and its runtime correction are coupled. Attention-temperature modulation (ATM) and output-head balancing (OHB) may compensate residual quantization drift, but applying either mechanism indiscriminately can perturb an already stable model; applying both can double-correct the same propagated error. A deployable method must therefore decide both *where* to quantize and *whether* a calibrated correction is warranted.

Existing mixed-precision methods rank layers using local reconstruction error, activation magnitude, Hessian sensitivity, or task-specific proxies Hubara et al. (2021); Frantar et al. (2023); Dong et al. (2019; 2020); Chen et al. (2025); Zhao et al. (2024). These signals are useful but incomplete for generative control. Centered kernel alignment (CKA) measures preservation of inter-sample geometry Raghu et al. (2017); Morcos et al. (2018); Kornblith et al. (2019), yet is intentionally insensitive to isotropic scaling. Cauchy–Schwarz (CS) divergence measures distribution overlap Yin et al. (2026), but does not say whether a shift changes the executed action. Neither local metric captures cancellation or amplification when many layers are quantized together, and neither determines whether a runtime correction should be admitted.

We address these challenges with GDSQ-VLA, a training-free W4/FP16 allocation and correction-selection framework. For each candidate Linear layer, an isolated W4A8 intervention is compared with a paired full-precision attribution path. CKA and CS provide complementary representation evidence; paired action-trajectory sensitivity weights that evidence; and activation-RMS and saturation guards reject infeasible assignments. A byte-constrained allocator generates diverse masks, after which complete-configuration functional adjudication measures cross-layer action error. The result is a frozen mixed-precision plan rather than a per-task policy.

The final method then applies a *calibration-gated runtime correction selector*. From model-level statistics of the frozen correction artifact, three preregistered thresholds either reject correction or admit the single dominant mechanism. Its output is restricted to the uncorrected baseline, ATM, or OHB; ATM+OHB is never an allowed final decision. The selector cannot inspect task identity, rollout success, or runtime feedback. The frozen v8 rule selects the baseline for GR00T N1.5 and OHB for $\pi_{0.5}$. On 256 frozen observations per model, each runtime path is bitwise identical to its corresponding static path, providing an explicit gate for any result reuse.

We evaluate on RoboCasa365 Nasiriany et al. (2026) with official horizons, four denoising steps, 16 executed actions per replan, and paired action noise. The completed GR00T selector result obtains 50.8% task-macro success over 50 tasks at a theoretical $1.99\times$ candidate-component compression, versus 55.1% for FP16. The original QuantVLA configuration obtains 30.4%, but uses a different compression rate and is therefore a secondary comparison rather than same-budget causal evidence. The final $\pi_{0.5}$ selector obtains 27.5% all-task macro SR. It improves over the uncorrected static mask (25.3%) by 2.2 points (95% task CI [0.6, 3.9], Holm-adjusted $p = 0.045$). Its 1.3-point difference from FP16 has CI [-0.5, 3.2] and Holm-adjusted $p = 0.240$, so this comparison is competitive rather than superior. Both uniform-W6 controls have complete audited coverage and are reported descriptively. The search-matched random and action-only controls remain excluded from superiority claims until their exact manifests reach full coverage.

Our contributions are threefold:

- (1) We formulate VLA precision allocation as action-weighted preservation of representation geometry and distribution overlap, with explicit stability guards and complete-mask functional adjudication under a frozen byte budget.
- (2) We introduce a model-level calibration gate that selects at most one runtime correction mechanism without task metadata or outcome feedback, and we verify its selected execution paths by bitwise equivalence.
- (3) We provide an auditable RoboCasa365 protocol in which manifests bind checkpoints, plans, code, environments, task/seed coverage, and paper cells; unsupported or incomplete results cannot populate the main table.

2 RELATED WORK

Efficient VLA inference. Generalist VLA policies combine large multimodal encoders with continuous-control heads Brohan et al. (2023a); Octo Model Team et al. (2024); Kim et al. (2024); Black et al. (2024); Physical Intelligence et al. (2025); NVIDIA et al. (2025). Efficiency methods reduce their cost through compact architectures, dynamic depth, token reduction, or temporal reuse Shukor et al. (2025); Zhang et al. (2025); Yang et al. (2025); Xu et al. (2025). SQAP-VLA co-designs quantization-aware token pruning with the quantizer to obtain structured inference acceleration Fang et al. (2025). These directions can reduce computation as well as storage, but modify

architecture, token flow, or execution. Our allocation leaves the checkpoint and operator graph intact; consequently, our present claim is static low-bit storage rather than latency acceleration.

VLA post-training quantization. QuantVLA adapts DuQuant-style transformations to language and diffusion-action stacks and calibrates attention temperature and output-head energy [Zhang et al. \(2026\)](#). ActQuant is the closest action-guided mixed-precision method: it assigns tensor bit widths from action contribution and optimizes block scales with action-aware curvature, targeting sub-4-bit deployment [Akbari et al. \(2026\)](#). Ω -QVLA instead pursues full-stack W4A4 quantization of both the language backbone and diffusion head using composite rotations and denoising-step scaling [Wang et al. \(2026\)](#). These methods establish that control signals and the action head require special treatment. GDSQ-VLA addresses a different question: which layers should retain FP16 under a fixed byte budget when local representation damage, downstream action sensitivity, and complete-mask interactions disagree. We preregister a pinned-code LIBERO reproduction and report only our locally audited values before making any shared-protocol ranking.

Transformer and diffusion PTQ. Weight-only methods such as GPTQ and AWQ optimize rounding or protect activation-salient channels [Frantar et al. \(2023\)](#); [Lin et al. \(2024b\)](#). Weight-activation PTQ reshapes difficult distributions through smoothing, calibration, or rotations [Xiao et al. \(2023\)](#); [Shao et al. \(2024\)](#); [Ashkboos et al. \(2024\)](#); [Lin et al. \(2024a\)](#); [Liu et al. \(2025b\)](#); [Sun et al. \(2025\)](#). Diffusion PTQ additionally handles timestep-dependent activation ranges and iterative error accumulation [Shang et al. \(2023\)](#); [Li et al. \(2023\)](#); [Wu et al. \(2024\)](#); [Chen et al. \(2025\)](#); [Zhao et al. \(2024\)](#). These techniques improve the low-bit numerical substrate. We reuse a DuQuant-style W4A8 substrate so that the studied variable is precision allocation and correction selection, not a new quantizer kernel.

Mixed precision and representation evidence. Mixed-precision allocation has been driven by hardware feedback, Hessian sensitivity, reconstruction loss, and task-specific diffusion metrics [Wang et al. \(2019\)](#); [Dong et al. \(2019; 2020\)](#); [Hubara et al. \(2021\)](#); [Li et al. \(2021\)](#); [Zhao et al. \(2024\)](#). SVCCA and projection-weighted CCA compare representation subspaces [Raghu et al. \(2017\)](#); [Morcos et al. \(2018\)](#), while CKA measures sample geometry with invariance to orthogonal transformations and isotropic scaling [Kornblith et al. \(2019\)](#). CS divergence provides a complementary distribution-overlap statistic [Yin et al. \(2026\)](#). GDSQ-VLA uses these metrics as frozen intervention probes, scales them with paired action sensitivity, and reserves action-functional evaluation for complete masks. Its runtime selector is separate from mask search: a preregistered model-level gate admits at most one calibrated correction and cannot use task identity or rollout outcomes.

3 METHOD

Overview. GDSQ-VLA converts a pretrained VLA policy into a frozen mixed-precision policy and a model-level correction decision without updating model parameters. The allocation pipeline in Fig. 1 has four stages: isolated layer interventions, dual-similarity and action scoring, constrained mask search, and paired functional adjudication. A fifth stage calibrates ATM and OHB for the selected mask and applies a preregistered absolute gate. The deployed output is therefore a pair (M, v) , where M is a W4/FP16 mask and $v \in \{\text{baseline}, \text{ATM}, \text{OHB}\}$ is fixed for the model.

3.1 ISOLATED QUANTIZATION INTERVENTIONS

Given visual observations I , instruction l , robot state s , and denoising time t , a VLA policy iteratively updates an action sample

$$x_{t-1} = f_{\theta}(x_t, F_{\text{VL}}(I, l), s, t). \quad (1)$$

We use group-wise DuQuant-style W4 weights and static A8 activations with group size 64 [Lin et al. \(2024a\)](#). The selectable inventory is checkpoint-specific. GR00T contains 116 candidate Linear layers—84 language projections and 32 DiT feed-forward projections—while 64 protected DiT attention projections remain FP16. The $\pi_{0.5}$ inventory contains 180 language and action-expert Linear layers.

A single reference path cannot separate W4 attribution from wrapper mechanics. We therefore use an *attribution reference* R that retains the transformations and A8 wrappers but routes every candidate weight through its full-precision branch. The *deployment reference* is the unwrapped FP16 policy.




Figure 1: GDSQ-VLA from frozen calibration to action generation. Single-layer interventions measure complementary geometric and distributional changes; action weighting, feasibility guards, a W6-equivalent byte budget, and paired functional adjudication freeze the W4/FP16 mask. A separate model-level calibration gate then selects exactly one of baseline, ATM, or OHB without task identity or rollout feedback.

For candidate i , all layers follow R except i , which uses W4. A fixed calibration buffer, token mask, and initial action noise are shared between the intervention and R . Hooks collect matched outputs (H_i^R, H_i^Q) ; padding rows are removed with the reference mask, and deterministic stratified pooling prevents visual, language, or state tokens from dominating. Representation and single-layer action scores use R , whereas activation guards and complete-mask adjudication use the deployment reference.

3.2 DUAL SIMILARITY, ACTION WEIGHTING, AND GUARDS

CKA measures preservation of inter-sample geometry. For centered matrices $\tilde{H}_i^R, \tilde{H}_i^Q \in \mathbb{R}^{N \times d}$, let $G_i^{UV} = (\tilde{H}_i^U)^\top \tilde{H}_i^V$. We compute

$$\text{CKA}_i = \frac{\|G_i^{RQ}\|_F^2}{\|G_i^{RR}\|_F \|G_i^{QQ}\|_F + \epsilon}, \quad D_i^{\text{CKA}} = 1 - \text{CKA}_i. \quad (2)$$

Its invariance to isotropic scaling is useful for stable geometry comparison but creates a scale blind spot. We complement it with an empirical CS divergence. For Gaussian kernel k whose bandwidth is fixed from the reference samples,

$$\begin{aligned} D_i^{\text{CS}} = & \log \left(\frac{1}{N^2} \sum_{a,b} k(h_a^R, h_b^R) \right) + \log \left(\frac{1}{N^2} \sum_{a,b} k(h_a^Q, h_b^Q) \right) \\ & - 2 \log \left(\frac{1}{N^2} \sum_{a,b} k(h_a^R, h_b^Q) \right). \end{aligned} \quad (3)$$

We use log-sum-exp accumulation and a bandwidth floor of 10^{-3} . After joint min-max normalization across candidates, the representation score is

$$S_i = \lambda_{\text{CKA}} \overline{D_i^{\text{CKA}}} + \lambda_{\text{CS}} \overline{D_i^{\text{CS}}}. \quad (4)$$

Representation damage matters only when it changes control. Under the same isolated intervention and paired denoising noise, let d_i denote action-trajectory divergence from R . We define

$$w_i = \frac{nd_i}{\sum_j d_j + \epsilon}, \quad (5)$$

then robustly clip and renormalize the weights so that $n^{-1} \sum_i w_i = 1$. This distinguishes layers with similar representation distortion but different downstream effects.

Similarity is not a stability certificate. Each intervention also records median log-RMS drift from unwrapped FP16 and the estimated fraction beyond the downstream static-A8 range. GR00T uses the frozen thresholds $\tau_{\text{RMS}} = 0.1319$ and $\tau_{\text{sat}} = 0.00232$; $\pi_{0.5}$ uses its separately calibrated thresholds. A violation forces FP16. The saturation statistic is an engineering approximation based on the calibrated consumer scale, not a direct measurement of fused-kernel clipping.

3.3 BYTE-CONSTRAINED SEARCH AND FUNCTIONAL ADJUDICATION

Let $x_i = 1$ assign W4A8 and $x_i = 0$ preserve the original FP16 Linear. Candidate masks minimize

$$\begin{aligned} \min_x \quad & \sum_i x_i w_i S_i, \\ \text{s.t.} \quad & \sum_i [x_i C_i^{\text{W4}} + (1 - x_i) C_i^{\text{FP16}}] \leq B, \\ & x_i = 0 \quad \text{for every guard violation,} \quad x_i \in \{0, 1\}. \end{aligned} \quad (6)$$

Here C_i is theoretical tightly packed weight storage. The uniform-W6 candidate-scope budgets are 862,912,512 bytes for GR00T and 2,042,542,080 bytes for $\pi_{0.5}$. They exclude activations, CUDA workspaces, simulator state, and fake-quant caches. Greedy allocation and local bit flips produce diverse low-objective masks; the additive optimum is not treated as the final policy.

Simultaneous interventions can cancel or amplify. We therefore instantiate selected Top- K masks with plan-specific A8 calibration and compare complete policies under paired observations and noise. The functional score is

$$D_{\text{func}} = D_{\text{final}} + D_{\text{kin}} + D_{\text{grip}} + 2 \text{CVaR}_{0.9}(D_{\text{solver}}), \quad (7)$$

where the tail term prevents a low mean from hiding severe denoising failures. Functional adjudication selects a representative for each CKA:CS ratio. On four predeclared GR00T development tasks, the frozen rule selects 16:1 from $\{8, 16, 32, 64, 128, 256\}$: 1 and a CKA-only boundary. The final masks assign W4 to 100 of 116 GR00T candidates and 80 of 180 $\pi_{0.5}$ candidates. A 10,000-draw task/seed bootstrap and leave-one-task-out analysis conducted after freezing finds no ratio whose paired interval against 16:1 excludes zero; we therefore interpret 16:1 as a selected point on a broad development plateau, not a universal optimum.

3.4 CALIBRATION-GATED RUNTIME CORRECTION

The mixed-precision mask can leave structured attention drift, but a correction can also destabilize an already adequate policy. For the frozen mask, ATM calibrates a query-temperature multiplier α and OHB calibrates per-head output multipliers β_h . We first choose one *candidate* mechanism from their model-level calibration geometry. OHB is the candidate when its correction is broad and its mean absolute log-magnitude is at least twice that of ATM. ATM is the candidate when corrected attention projections remain FP16 and its directional coherence exceeds OHB by at least 0.05. Otherwise the candidate is the baseline. ATM and OHB are not composed because their ratios act at different points of the same propagated attention error and are not calibrated as an additive correction.

The v8 absolute gate then decides whether to admit the candidate. For correction values $c_1, \dots, c_m$, define

$$\begin{aligned} q_{\text{dir}} &= \frac{|n_{<1} - n_{>1}|}{m}, & q_{\text{non}} &= \frac{n_{c \neq 1}}{m}, \\ q_{\text{log}} &= \frac{1}{m} \sum_j |\log c_j|. \end{aligned} \quad (8)$$

The correction is enabled only if $q_{\text{dir}} \geq 0.30$, $q_{\text{non}} \geq 0.85$, and $q_{\text{log}} \geq 0.05$; otherwise the baseline is selected. These three thresholds and the selector SHA were frozen before formal testing. GR00T's ATM candidate fails all three gates and therefore selects the uncorrected GDSQ mask. The $\pi_{0.5}$ OHB candidate passes all three and selects OHB. The decision is evaluated once per model, so task names retained in runtime metadata are attestations only and cannot affect the selected variant.

Every formal result row records the selector SHA, rule name, selected variant, selected configuration ID, and server-metadata hash. For GR00T, whose bitwise gate permits reuse of the immutable static matrix, a separate 2,500-row derived attestation adds these selector fields and the source-row hash without modifying the original outcomes. The runtime rejects an unknown variant and v8 rejects ATM+OHB. Consequently, the final method is exactly GDSQ-VLA plus one frozen model-level choice from {baseline, ATM, OHB}; static GDSQ, ATM, OHB, and ATM+OHB configurations are ablations.

4 EXPERIMENTS

4.1 EXPERIMENTAL SETUP

Benchmark and policies. We evaluate GR00T N1.5 and $\pi_{0.5}$ [Physical Intelligence et al. \(2025\)](#) on the RoboCasa365 v1.0.1 target split [Nasiriany et al. \(2026\)](#). The benchmark has 18 Atomic-Seen, 16 Composite-Seen, and 16 Composite-Unseen tasks. GR00T uses the corresponding 60k-step post-training checkpoints; $\pi_{0.5}$ uses its RoboCasa checkpoint. Four atomic tasks—OpenCabinet, OpenStandMixerHead, PickPlaceDrawerToCounter, and CoffeeSetupMug—form the predeclared GR00T ratio-development set. The $\pi_{0.5}$ allocation transfers the frozen 16:1 ratio and performs model-specific functional adjudication and calibration; it does not tune the ratio from $\pi_{0.5}$ test rollouts.

Configurations and budgets. The final configuration is GDSQ-VLA plus the frozen v8 runtime selector. The selector chooses the uncorrected GDSQ mask for GR00T and OHB for $\pi_{0.5}$. Static GDSQ, ATM, OHB, and ATM+OHB are ablations, not alternative definitions of ours. FP16 is the uncompressed reference. QuantVLA W4A8 [Zhang et al. \(2026\)](#) is retained as a secondary low-bit comparison because its compression rate differs from ours. The direct budget controls are uniform W6, a search-matched random W4/FP16 mask, and an action-only allocator. Random search uses the same 26 candidates, functional-evaluation count, and seven development representatives as GDSQ; its seeds and masks are frozen before held-out rollouts. All masks obey the same uniform-W6 candidate-byte ceiling and use the same checkpoint, A8 calibration strategy, and guards.

Rollout and statistics. Every formal configuration uses 50 seeds per task, official horizons, a fresh rendered environment, four denoising steps, and 16 executed actions per replan. Initial action noise is deterministically keyed by (task, seed, replan), giving paired policy comparisons without changing the marginal normal distribution. We report task-macro SR, 10,000-draw task-cluster bootstrap intervals, paired task-level sign-flip tests, and Holm correction over preregistered contrasts. A no-guards crash, NaN, or out-of-range action is a formal failure rather than a silently dropped episode.

Claim–evidence gate. Each run manifest binds checkpoint, plan, calibration, selector, launcher, evaluator, environment, task list, and seed coverage by SHA-256. Each result row records its server-metadata hash and, when applicable, selector SHA, rule, variant, and configuration ID. The main table is generated from the evidence registry; a task-set cell remains pending unless its declared coverage is exact, duplicates are absent, statistics are reproducible, and source drift has an explicit audit. Replay audits retain the frozen launch manifests: they reconstruct every task/seed key and must reproduce the registered summary exactly before later source edits can be disclosed without invalidating immutable result rows. This gate is passed by the completed $\pi_{0.5}$ four-configuration matrix, the three-manifest GR00T uniform-W6 evaluation, and all three Ω -QVLA task sets for both policy families.

External Table 1 reproduction. For Ω -QVLA [Wang et al. \(2026\)](#), we pin the authors’ code commit but rebuild model- and task-set-specific W4A4 packs from frozen RoboCasa365 calibration data because the released LIBERO packs cannot be transferred. Table 1 reports complete 50-task evaluations for both GR00T and $\pi_{0.5}$. Table 2 separately uses the same five configurations as Table 1—FP16, QuantVLA W4A8, Uniform W6, Ω -QVLA W4A4, and ours—under the original LIBERO protocol: four suites, ten tasks per suite, ten held-out trials per task, and eight denoising steps. We import no author-reported values; all 4,000 local episodes must pass the joint manifest gate.

Table 1: RoboCasa365 task-macro success rate (SR, %) over 50 paired scenarios per task.

Configuration	Low-bit layers	Atomic SR $\uparrow$	C-Seen SR $\uparrow$	C-Unseen SR $\uparrow$	All SR $\uparrow$	Size (GiB) $\downarrow$	Comp. $\uparrow$
GR00T N1.5 (<i>target split</i>)							
FP16	—	75.6	41.9	45.3	55.1	1.993	1.00 $\times$
QUANTVLA W4A8	116 W4	49.6	19.6	19.8	30.4	0.898	2.22 $\times$
Uniform W6	116 W6	68.4	42.9	41.8	51.7	1.109	1.80 $\times$
Ω -QVLA W4A4 [‡]	180 W4	60.1	25.9	27.5	38.7	0.599	3.33 $\times$
GDSQ-VLA(Ours)	100 W4	69.0	40.6	40.4	50.8	1.001	1.99 $\times$
$\pi_{0.5}$ Physical Intelligence et al. (2025) (<i>target split</i>)							
FP16	—	58.3	14.0	2.1	26.2	4.113	1.00 $\times$
QUANTVLA W4A8	180 W4	56.1	12.8	1.8	24.8	1.388	2.96 $\times$
Uniform W6	180 W6	56.6	10.9	2.1	24.5	1.902	2.16 $\times$
Ω -QVLA W4A4 [‡]	252 W4	49.6	10.4	1.0	21.5	1.307	3.27 $\times$
GDSQ-VLA(Ours)	80 W4	57.8	17.0	3.9	27.5	1.868	2.20 $\times$

C-Seen/C-Unseen denote Composite-Seen/Composite-Unseen. Ours uses the frozen v8 selector. Uniform W6 is the direct byte-budget baseline; QuantVLA is a different-rate comparison. Size is theoretical tightly packed candidate-Linear storage. [‡] Ω -QVLA uses model- and task-set-specific RoboCasa365 calibration; both rows have exact 50-task $\times$ 50-scenario coverage. Their sizes include packed W4 weights and scales/rotation metadata; the $\pi_{0.5}$ compression ratio uses its audited 252-layer FP16 scope.

4.2 SELECTOR VALIDATION

The path-equivalence test uses 256 frozen observations per model, four denoising steps, and paired noise. GR00T compares the selector’s baseline decision with static GDSQ; $\pi_{0.5}$ compares the selector’s OHB decision with static OHB. Both comparisons are bitwise equal on all 256/256 action tensors, with maximum absolute difference zero and no non-finite output. This passes the preregistered reuse gate for the completed GR00T static 50-task run. It does *not* authorize reuse of a partial $\pi_{0.5}$ OHB diagnostic, so we run a separate formal 50-task selector evaluation for $\pi_{0.5}$.

Before formal evaluation, a two-task, two-seed closed-loop preflight verifies four complete episodes and all row-level attestations; this preflight is diagnostic only. The $\pi_{0.5}$ manifest is then frozen before the first official rollout and specifies exactly 2,500 unique task/seed keys. The audited run completes all 2,500 manifested episodes; its three task-group macro SRs are 57.8%, 17.0%, 3.9%, for 27.5% over all 50 tasks. All 2,500 rows select OHB and attest the frozen selector SHA; no task metadata or runtime outcome enters the decision.

4.3 COMPLETED ROBOCASA365 EVIDENCE

GR00T final selector. By bitwise reuse, the final GR00T selector obtains 69.0%, 40.6%, and 40.4% task-macro SR on Atomic, Composite-Seen, and Composite-Unseen, respectively, for 50.8% over all 50 tasks. FP16 obtains 55.1%; the paired all-task difference for ours minus FP16 is -4.32 points with a 95% task-cluster interval of $[-6.56, -2.08]$. The QuantVLA configuration obtains 30.4%, and ours exceeds it by 20.3 points with interval $[16.9, 24.0]$. Because QuantVLA compresses the studied scope by 2.22 $\times$ while ours compresses it by 1.99 $\times$, this is a descriptive secondary comparison, not evidence that the selection signal is superior at fixed storage.

The completed uniform-W6 control obtains 68.4%, 42.9%, and 41.8% on the three task groups, for 51.7% overall at 1.80 $\times$ theoretical compression. Ours is -1.0 point relative to W6, with a paired 95% task-cluster interval of $[-3.1, 1.2]$, while using 1.001 rather than 1.109 GiB in the candidate Linear scope. We therefore treat the GR00T fixed-ceiling comparison as competitive storage–accuracy evidence, not superiority.

Ω -QVLA RoboCasa365 reproduction. The completed GR00T evaluation contains exactly 2,500 episodes—50 tasks with 50 paired seeds each—and no formal failure. Ω -QVLA obtains 60.1%, 25.9%, and 27.5% task-macro SR on Atomic, Composite-Seen, and Composite-Unseen, respectively, for 38.7% over all tasks. The corresponding task-cluster intervals are $[50.7, 68.7]$, $[15.8, 36.6]$,

[16.0, 40.4], and [31.4, 46.3]. Thus W4A4 retains substantially more success on Atomic tasks than on the two compositional groups in this reproduction. Its theoretical tightly packed Linear-component size is 0.599 GiB, or $3.33\times$ relative to the same 1.993-GiB FP16 scope. Because this rate differs from ours and uniform W6, we report the result descriptively rather than as a fixed-budget ranking.

The independent $\pi_{0.5}$ reproduction likewise contains exactly 2,500 episodes with no missing, duplicate, or failed row. It obtains 49.6% Atomic SR (10,000-draw task-cluster 95% CI [37.0, 61.9]), 10.4% Composite-Seen SR ([4.6, 16.8]), and 1.0% Composite-Unseen SR ([0.1, 2.1]), for 21.5% over all 50 tasks ([14.1, 29.6]). The sharp compositional drop, especially on unseen tasks, exposes a policy-family boundary for full-stack W4A4 in this reproduction. Its audited 252-layer packed estimate is 1.307 GiB, or $3.27\times$ relative to its 4.271-GiB FP16 scope. Because this rate differs from ours and uniform W6, we again treat it as descriptive rather than a fixed-budget ranking.

$\pi_{0.5}$ final selector and architecture boundary. The final $\pi_{0.5}$ selector obtains 27.5% all-task macro SR. It improves over the uncorrected static mask (25.3%) by 2.2 points (95% task CI [0.6, 3.9], Holm-adjusted $p = 0.045$). Its 1.3-point difference from FP16 has CI [-0.5, 3.2] and Holm-adjusted $p = 0.240$, so this comparison is competitive rather than superior. The selector also exceeds QuantVLA W4A8 by 2.6 points (95% task CI [0.9, 4.5], Holm-adjusted $p = 0.023$), but QuantVLA has a different $2.96\times$ compression rate and remains a secondary comparison. The three task-group SRs of the final selector are 57.8%, 17.0%, and 3.9%, yielding 27.5% over all 50 tasks at $2.20\times$ theoretical candidate-component compression.

The uncorrected static mask obtains 25.3%, while static ATM+OHB reaches 26.6% as an ablation. The selector’s significant gain over the same frozen mask isolates the value of admitting OHB for this model; it does not establish that the allocation is universally superior. The completed uniform-W6 control obtains 24.5% at $2.16\times$ theoretical compression, giving a descriptive ours-minus-W6 gap of 3.0 points at nearly matched storage. We do not interpret this gap as superiority before the remaining same-budget controls and preregistered joint Holm correction are complete. In particular, the final selector’s 1.3-point difference from FP16 has an interval crossing zero. We therefore report $\pi_{0.5}$ as an architecture boundary: runtime correction helps the transferred allocation, while fixed-budget superiority remains unproven.

Same-budget causal controls. Uniform W6, shown in both main tables, isolates whether mixed precision helps under the same byte ceiling. Random masks control candidate count, functional-search budget, and retained FP16 bytes. The action-only allocator tests whether paired behavioral sensitivity alone explains the selected layers. The result-blind functional screens for both search controls are complete, but their closed-loop development selection and formal held-out rollouts are not. These two search-mechanism controls remain supplemental because they answer causal questions rather than define deployable headline configurations; they will be reported only after complete audited manifests. Superiority wording additionally requires the paired 95% interval for ours minus the strongest same-budget control to exclude zero under the preregistered Holm correction. Both uniform-W6 controls have complete audited coverage and are reported descriptively. The search-matched random and action-only controls remain excluded from superiority claims until their exact manifests reach full coverage.

4.4 LIBERO COMPARISON

The reproduction uses the pinned release and official quantization packs for Ω -QVLA and independently calibrates QuantVLA, Uniform W6, and GDSQ-VLA on the original GR00T N1.5 and $\pi_{0.5}$ checkpoints. We freeze the commands before observing local rollouts, use initialization offset 10 for the ten held-out states, and recalibrate the GDSQ selector per model without reusing RoboCasa artifacts. The audit preserves two release discrepancies: the paper and current builder documentation use different GPTQ damping values, and the released $\pi_{0.5}$ pack has ten per-step buckets while the evaluation path uses eight denoising steps. We load the released packs as-is without rebuilding or test-set retuning.

Table 2: LIBERO local comparison under the Ω -QVLA Table 1 protocol (success rate, %).

Configuration	Goal $\uparrow$	Spatial $\uparrow$	Object $\uparrow$	Long $\uparrow$	Avg. $\uparrow$
GR00T N1.5					
FP16	<i>pending</i>	<i>pending</i>	<i>pending</i>	<i>pending</i>	<i>pending</i>
QUANTVLA W4A8	<i>pending</i>	<i>pending</i>	<i>pending</i>	<i>pending</i>	<i>pending</i>
Uniform W6	<i>pending</i>	<i>pending</i>	<i>pending</i>	<i>pending</i>	<i>pending</i>
Ω -QVLA W4A4	<i>pending</i>	<i>pending</i>	<i>pending</i>	<i>pending</i>	<i>pending</i>
GDSQ-VLA(Ours)	<i>pending</i>	<i>pending</i>	<i>pending</i>	<i>pending</i>	<i>pending</i>
$\pi_{0.5}$					
FP16	<i>pending</i>	<i>pending</i>	<i>pending</i>	<i>pending</i>	<i>pending</i>
QUANTVLA W4A8	<i>pending</i>	<i>pending</i>	<i>pending</i>	<i>pending</i>	<i>pending</i>
Uniform W6	<i>pending</i>	<i>pending</i>	<i>pending</i>	<i>pending</i>	<i>pending</i>
Ω -QVLA W4A4	<i>pending</i>	<i>pending</i>	<i>pending</i>	<i>pending</i>	<i>pending</i>
GDSQ-VLA(Ours)	<i>pending</i>	<i>pending</i>	<i>pending</i>	<i>pending</i>	<i>pending</i>

Four official suites, ten tasks per suite, and ten held-out trials per task. All cells remain pending until the joint five-configuration manifest reaches exactly 4,000 valid local episodes; no author-reported values are imported. The official Ω -QVLA W4A4 packs are pinned. QuantVLA, Uniform W6, and GDSQ-VLA are independently calibrated on LIBERO; the GDSQ selector is model-specific, with no RoboCasa artifact reuse or test-set retuning.

4.5 RATIO STABILITY AND COMPONENT TESTS

The frozen point rule selects CKA:CS= 16:1 at 80.5% development SR. Post-freeze hierarchical task/seed bootstrap gives it the largest fractional best probability (0.486), but every alternative’s paired interval against 16:1 includes zero. Leave-one-task-out selects 16:1 on three omissions and ties it with 256:1 on the fourth. We therefore describe a broad ratio plateau rather than sharp 16:1 optimality. The ratio table, selection-controlled Primary-14 comparison, and earlier five-seed signal diagnostic are moved to the supplement. The new 50-seed component matrix—CKA-only, CS-only, $w_i = 1$, no guards, and no functional adjudication—is preregistered on GR00T Primary-14; incomplete cells cannot replace those older diagnostics.

Storage and runtime scope. Table 1 reports theoretical tightly packed candidate-Linear storage, not live GPU memory or latency. The present eager fake-quant path retains floating-point weights and GEMMs; the appendix details its accounting and runtime diagnostic.

5 CONCLUSION AND LIMITATIONS

This paper presents GDSQ-VLA, a training-free framework that couples action-weighted geometry/distribution preservation with byte-constrained W4/FP16 allocation and a calibration-gated selector over baseline, ATM, and OHB. The selector uses neither task identity nor outcome feedback, and its chosen paths are bitwise exact on 256 frozen observations per policy family. GR00T obtains 50.8% task-macro SR at $1.99\times$ theoretical candidate-component compression, versus 55.1% for FP16. The final $\pi_{0.5}$ selector obtains 27.5% all-task macro SR. It improves over the uncorrected static mask (25.3%) by 2.2 points (95% task CI [0.6, 3.9], Holm-adjusted $p = 0.045$). Its 1.3-point difference from FP16 has CI [-0.5, 3.2] and Holm-adjusted $p = 0.240$, so this comparison is competitive rather than superior.

The evidence also defines clear boundaries. Differences from QuantVLA use different compression rates; Same-budget controls must complete before supporting a causal superiority claim. The 16:1 ratio lies on a broad four-task development plateau and is transferred to $\pi_{0.5}$. The eager fake-quant path demonstrates theoretical static storage only, not live-memory or latency gains.

Future work will complete LIBERO and same-budget controls, calibration and storage–SR robustness, and packed-kernel memory/latency measurements. Shared-protocol external comparisons require reproducible releases; fused INT4/INT8 kernels and real-robot evaluation remain necessary for an end-to-end deployment claim.

AI USE STATEMENT

Generative AI tools were used to assist experiment planning, software implementation and debugging, result aggregation and interpretation, drafting and editing, reference and table formatting, and creation and post-processing of the architecture figure. The authors reviewed the generated material, tested AI-assisted code, reran the evidence and document gates, verified numerical claims against frozen manifests and audited summaries, and checked references against primary sources. No experimental result was synthesized by a generative model. The authors take responsibility for the final content, including all AI-assisted text, code, claims, and artifacts.

REPRODUCIBILITY STATEMENT

Sections 3 and 4 specify the allocation, selector, calibration, evaluation, and statistical protocols. The appendix reports the frozen ratio analysis, controlled comparison, component ablations, and correction diagnostics. The supplementary artifact includes versioned manifests, configuration and code hashes, exact task/seed coverage, table generators, and claim–evidence audits; every numerical paper cell is generated from a frozen summary rather than entered by hand.

REFERENCES

- Arash Akbari, Arman Akbari, Masih Eskandar, Qitao Tan, Yixiao Chen, Jingwu Luo, Bertha Pangaribuan, Liyun Zhang, Jennifer Dy, Geng Yuan, Xue Lin, Gaowen Liu, Stratis Ioannidis, and Yanzhi Wang. Actquant: Sub-4-bit action-guided quantization for vision-language-action models. *arXiv preprint arXiv:2605.24011*, 2026. URL <https://arxiv.org/abs/2605.24011>.
- Saleh Ashkboos, Amirkeivan Mohtashami, Maximilian Croci, Bo Li, Martin Jaggi, Dan Alistarh, Torsten Hoeffler, and James Hensman. Quarot: Outlier-free 4-bit inference in rotated llms. In *Advances in Neural Information Processing Systems*, 2024.
- Kevin Black, Noah Brown, Danny Driess, Adnan Esmail, Michael Equi, Chelsea Finn, Niccolo Fusai, Lachy Groom, Karol Hausman, Brian Ichter, et al. π_0 : A vision-language-action flow model for general robot control. *arXiv preprint arXiv:2410.24164*, 2024. URL <https://arxiv.org/abs/2410.24164>.
- Anthony Brohan, Noah Brown, Justice Carbajal, Yevgen Chebotar, Xi Chen, Krzysztof Choromanski, Tianli Ding, Danny Driess, Avinava Dubey, Chelsea Finn, et al. Rt-2: Vision-language-action models transfer web knowledge to robotic control. *Conference on Robot Learning*, 2023a. URL <https://arxiv.org/abs/2307.15818>.
- Anthony Brohan, Noah Brown, Justice Carbajal, Yevgen Chebotar, Joseph Dabis, Chelsea Finn, Keerthana Gopalakrishnan, Karol Hausman, Alex Herzog, Jasmine Hsu, et al. Rt-1: Robotics transformer for real-world control at scale. *arXiv preprint arXiv:2212.06817*, 2023b. URL <https://arxiv.org/abs/2212.06817>.
- Lei Chen, Yuan Meng, Chen Tang, Xinzhu Ma, Jingyan Jiang, Xin Wang, Zhi Wang, and Wenwu Zhu. Q-dit: Accurate post-training quantization for diffusion transformers. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 28306–28315, 2025. doi: 10.1109/CVPR52734.2025.02636.
- Cheng Chi, Zhenjia Xu, Siyuan Feng, Eric Cousineau, Yilun Du, Benjamin Burchfiel, Russ Tedrake, and Shuran Song. Diffusion policy: Visuomotor policy learning via action diffusion. In *Robotics: Science and Systems*, 2023. doi: 10.15607/RSS.2023.XIX.026.
- Zhen Dong, Zhewei Yao, Amir Gholami, Michael Mahoney, and Kurt Keutzer. Hawq: Hessian aware quantization of neural networks with mixed-precision. In *IEEE/CVF International Conference on Computer Vision*, 2019. URL <https://arxiv.org/abs/1905.03696>.
- Zhen Dong, Zhewei Yao, Yaohui Cai, Daiyaan Arfeen, Amir Gholami, Michael W. Mahoney, and Kurt Keutzer. Hawq-v2: Hessian aware trace-weighted quantization of neural networks. In *Advances in Neural Information Processing Systems*, 2020.

- Danny Driess, Fei Xia, Mehdi S. M. Sajjadi, Corey Lynch, Aakanksha Chowdhery, Brian Ichter, Ayaan Wahid, Jonathan Tompson, Quan Vuong, Tianhe Yu, et al. Palm-e: An embodied multimodal language model. In *International Conference on Machine Learning*, 2023. URL <https://arxiv.org/abs/2303.03378>.
- Hengyu Fang, Yijiang Liu, Yuan Du, Li Du, and Huanrui Yang. Sqap-vla: A synergistic quantization-aware pruning framework for high-performance vision-language-action models. *arXiv preprint arXiv:2509.09090*, 2025. URL <https://arxiv.org/abs/2509.09090>.
- Elias Frantar, Saleh Ashkboos, Torsten Hoefer, and Dan Alistarh. Gptq: Accurate post-training quantization for generative pre-trained transformers. In *International Conference on Learning Representations*, 2023. URL <https://arxiv.org/abs/2210.17323>.
- Itay Hubara, Yury Nahshan, Yair Hanani, Ron Banner, and Daniel Soudry. Accurate post training quantization with small calibration sets. In *International Conference on Machine Learning*, volume 139, pp. 4466–4475. PMLR, 2021. URL <https://proceedings.mlr.press/v139/hubara21a.html>.
- Moo Jin Kim, Karl Pertsch, Siddharth Karamcheti, Ted Xiao, Ashwin Balakrishna, Suraj Nair, Rafael Rafailov, Ethan Foster, Pannag Sanketi, Quan Vuong, et al. Openvla: An open-source vision-language-action model. *arXiv preprint arXiv:2406.09246*, 2024. URL <https://arxiv.org/abs/2406.09246>.
- Simon Kornblith, Mohammad Norouzi, Honglak Lee, and Geoffrey Hinton. Similarity of neural network representations revisited. In *International Conference on Machine Learning*, 2019. URL <https://arxiv.org/abs/1905.00414>.
- Xiuyu Li, Yijiang Liu, Long Lian, Huanrui Yang, Zhen Dong, Daniel Kang, Shanghang Zhang, and Kurt Keutzer. Q-diffusion: Quantizing diffusion models. In *IEEE/CVF International Conference on Computer Vision*, pp. 17489–17499, 2023. doi: 10.1109/ICCV51070.2023.01608.
- Yuhang Li, Ruihao Gong, Xu Tan, Yang Yang, Peng Hu, Qi Zhang, Fengwei Yu, Wei Wang, and Shi Gu. Brecq: Pushing the limit of post-training quantization by block reconstruction. In *International Conference on Learning Representations*, 2021. URL <https://arxiv.org/abs/2102.05426>.
- Haokun Lin, Haobo Xu, Yichen Wu, Jingzhi Cui, Yingtao Zhang, Linzhan Mou, Linqi Song, Zhenan Sun, and Ying Wei. Duquant: Distributing outliers via dual transformation makes stronger quantized llms. In *Advances in Neural Information Processing Systems*, 2024a.
- Ji Lin, Jiaming Tang, Haotian Tang, Shang Yang, Wei-Ming Chen, Wei-Chen Wang, Guangxuan Xiao, Xingyu Dang, Chuang Gan, and Song Han. Awq: Activation-aware weight quantization for llm compression and acceleration. *Proceedings of Machine Learning and Systems*, 6:87–100, 2024b.
- Songming Liu, Lingxuan Wu, Bangguo Li, Hengkai Tan, Huayu Chen, Zhengyi Wang, Ke Xu, Hang Su, and Jun Zhu. Rdt-1b: A diffusion foundation model for bimanual manipulation. *arXiv preprint arXiv:2410.07864*, 2025a. URL <https://arxiv.org/abs/2410.07864>.
- Zechun Liu, Changsheng Zhao, Igor Fedorov, Bilge Soran, Dhruv Choudhary, Raghuraman Krishnamoorthi, Vikas Chandra, Yuandong Tian, and Tijmen Blankevoort. Spingquant: Llm quantization with learned rotations. In *International Conference on Learning Representations*, 2025b. URL <https://arxiv.org/abs/2405.16406>.
- Ari S. Morcos, Maithra Raghu, and Samy Bengio. Insights on representational similarity in neural networks with canonical correlation. In *Advances in Neural Information Processing Systems*, 2018.
- Soroush Nasiriany, Sepehr Nasiriany, Abhiram Maddukuri, and Yuke Zhu. Robocasa365: A large-scale simulation framework for training and benchmarking generalist robots. In *International Conference on Learning Representations*, 2026. URL <https://arxiv.org/abs/2603.04356>.

- NVIDIA, Johan Bjorck, Fernando Castañeda, Nikita Cherniadev, Xingye Da, Runyu Ding, Linxi Fan, et al. Gr00t n1: An open foundation model for generalist humanoid robots. *arXiv preprint arXiv:2503.14734*, 2025. URL <https://arxiv.org/abs/2503.14734>.
- Octo Model Team, Dibya Ghosh, Homer Walke, Karl Pertsch, Kevin Black, Oier Mees, Sudeep Dasari, Joey Hejna, Tobias Kreiman, Charles Xu, et al. Octo: An open-source generalist robot policy. *arXiv preprint arXiv:2405.12213*, 2024. URL <https://arxiv.org/abs/2405.12213>.
- Physical Intelligence, Kevin Black, Noah Brown, James Darpinian, Karan Dhabalia, Danny Driess, Adnan Esmail, Michael Equi, Chelsea Finn, Niccolo Fusai, Manuel Y. Galliker, Dibya Ghosh, Lachy Groom, Karol Hausman, Brian Ichter, Szymon Jakubczak, Tim Jones, Liyiming Ke, Devin LeBlanc, Sergey Levine, Adrian Li-Bell, Mohith Mothukuri, Suraj Nair, Karl Pertsch, Allen Z. Ren, Lucy Xiaoyang Shi, Laura Smith, Jost Tobias Springenberg, Kyle Stachowicz, James Tanner, Quan Vuong, Homer Walke, Anna Walling, Haohuan Wang, Lili Yu, and Ury Zhilinsky. $\pi_{0.5}$: A vision-language-action model with open-world generalization. *arXiv preprint arXiv:2504.16054*, 2025. doi: 10.48550/arXiv.2504.16054. URL <https://arxiv.org/abs/2504.16054>.
- Maithra Raghu, Justin Gilmer, Jason Yosinski, and Jascha Sohl-Dickstein. Svcca: Singular vector canonical correlation analysis for deep learning dynamics and interpretability. In *Advances in Neural Information Processing Systems*, 2017.
- Yuzhang Shang, Zhihang Yuan, Bin Xie, Bingzhe Wu, and Yan Yan. Post-training quantization on diffusion models. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 1972–1981, 2023. doi: 10.1109/CVPR52729.2023.00196.
- Wenqi Shao, Mengzhao Chen, Zhaoyang Zhang, Peng Xu, Lirui Zhao, Zhiqian Li, Kaipeng Zhang, Peng Gao, Yu Qiao, and Ping Luo. Omniquant: Omnidirectionally calibrated quantization for large language models. In *International Conference on Learning Representations*, 2024. URL <https://arxiv.org/abs/2308.13137>.
- Mustafa Shukor, Dana Aubakirova, Francesco Capuano, Pepijn Kooijmans, Steven Palma, Adil Zouitine, Michel Aractingi, Caroline Pascal, Martino Russi, Andres Marafioti, et al. Smolvla: A vision-language-action model for affordable and efficient robotics. *arXiv preprint arXiv:2506.01844*, 2025. URL <https://arxiv.org/abs/2506.01844>.
- Yuxuan Sun, Ruikang Liu, Haoli Bai, Han Bao, Kang Zhao, Yuening Li, Jiaxin Hu, Xianzhi Yu, Lu Hou, Chun Yuan, et al. Flatquant: Flatness matters for llm quantization. In *International Conference on Machine Learning*, 2025. URL <https://arxiv.org/abs/2410.09426>.
- Kuan Wang, Zhijian Liu, Yujun Lin, Ji Lin, and Song Han. Haq: Hardware-aware automated quantization with mixed precision. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2019.
- Xinyu Wang, Mingze Li, Sicheng Lyu, Dongxiu Liu, Kaicheng Yang, Ziyu Zhao, Yufei Cui, Xiao-Wen Chang, and Peng Lu. Omega-QVLA: Robust quantization for vision-language-action models via composite rotation and per-step scaling. *arXiv preprint arXiv:2605.28803*, 2026. URL <https://arxiv.org/abs/2605.28803v1>.
- Junyi Wu, Haoxuan Wang, Yuzhang Shang, Mubarak Shah, and Yan Yan. Ptq4dit: Post-training quantization for diffusion transformers. In *Advances in Neural Information Processing Systems*, volume 37, pp. 62732–62755, 2024. doi: 10.52202/079017-2006.
- Guangxuan Xiao, Ji Lin, Mickael Seznec, Hao Wu, Julien Demouth, and Song Han. Smoothquant: Accurate and efficient post-training quantization for large language models. In *International Conference on Machine Learning*, 2023.
- Siyu Xu, Yunke Wang, Chenghao Xia, Dihao Zhu, Tao Huang, and Chang Xu. Vla-cache: Efficient vision-language-action manipulation via adaptive token caching. In *Advances in Neural Information Processing Systems*, 2025. URL <https://arxiv.org/abs/2502.02175>.
- Yantai Yang, Yuhao Wang, Zichen Wen, Luo Zhongwei, Chang Zou, Zhipeng Zhang, Chuan Wen, and Linfeng Zhang. Efficientvla: Training-free acceleration and compression for vision-language-action models. *arXiv preprint arXiv:2506.10100*, 2025. URL <https://arxiv.org/abs/2506.10100>.

- Wenzhe Yin, Zehao Xiao, Pan Zhou, Shujian Yu, Jiayi Shen, Jan-Jakob Sonke, and Efstratios Gavves. Distributional vision-language alignment by cauchy–schwarz divergence. In *International Conference on Learning Representations*, 2026. URL <https://arxiv.org/abs/2502.17028>.
- Jingxuan Zhang, Yunta Hsieh, Zhongwei Wan, Haokun Lin, Xin Wang, Ziqi Wang, Yingtie Lei, and Mi Zhang. Quantvla: Scale-calibrated post-training quantization for vision-language-action models. *arXiv preprint arXiv:2602.20309*, 2026. URL <https://arxiv.org/abs/2602.20309>.
- Rongyu Zhang, Menghang Dong, Yuan Zhang, Liang Heng, Xiaowei Chi, Gaole Dai, Li Du, Yuan Du, and Shanghang Zhang. Mole-vla: Dynamic layer-skipping vision language action model via mixture-of-layers for efficient robot manipulation. *arXiv preprint arXiv:2503.20384*, 2025. URL <https://arxiv.org/abs/2503.20384>.
- Tianchen Zhao, Xuefei Ning, Tongcheng Fang, Enshu Liu, Guyue Huang, Zinan Lin, Shengen Yan, Guohao Dai, and Yu Wang. Mixdq: Memory-efficient few-step text-to-image diffusion models with metric-decoupled mixed precision quantization. In *European Conference on Computer Vision*, pp. 285–302, 2024. doi: 10.1007/978-3-031-72630-9_17.

Table 3: Predeclared ratio selection on four development tasks (50 scenarios per task). Lower D_{func} is better, but the final ratio is selected by closed-loop macro success rate.

CKA:CS ratio	$D_{\text{func}} \downarrow$	Dev SR (%) $\uparrow$
8:1	0.11597	77.0
16:1	0.10044	80.5
32:1	0.13395	76.5
64:1	0.07495	75.0
128:1	0.10236	77.0
256:1	0.09328	79.5
CKA only	0.09157	77.5

Table 4: Selection-controlled GR00T N1.5 Atomic results. The CKA:CS ratio is frozen using four development tasks; none of the Primary 14 tasks contributes to ratio selection. The final selector chooses the baseline path and is bitwise identical to the static GDSQ-VLA row. Confidence intervals are 95% task-cluster bootstrap intervals.

Configuration	Primary 14 SR (%) $\uparrow$
FP16 (uncompressed)	73.9 [66.0, 80.9]
QUANTVLA W4A8	45.6 [34.4, 57.0]
GDSQ-VLA + ATM/OHB	63.0 [52.0, 72.9]
GDSQ-VLA + selector	65.9 [57.3, 73.7]

A SUPPLEMENTARY EVIDENCE

A.1 RATIO SWEEP AND STABILITY

The point-selection rule freezes 16:1 because it has the largest development task-macro SR. A separate post-freeze analysis resamples tasks and then paired seeds for 10,000 draws. No alternative ratio has a paired difference interval against 16:1 that excludes zero: for example, the 256:1 interval is $[-6.5, 5.0]$ points and the CKA-only interval is $[-10.0, 3.5]$ points. Leave-one-task-out chooses 16:1 for three omissions and ties 16:1 with 256:1 for the fourth. These diagnostics support stability within a plateau, not universal ratio optimality.

A.2 SELECTION-CONTROLLED GR00T ATOMIC RESULT

The four ratio-development tasks are excluded from this comparison. The final GR00T selector chooses the baseline path and is bitwise identical to the static GDSQ-VLA row. It obtains 65.9% SR on Primary-14, versus 45.6% for QuantVLA W4A8; the paired difference is +20.3 points with interval $[13.9, 26.9]$. FP16 obtains 73.9%. Because QuantVLA and GDSQ-VLA use different static byte totals, the result is a selection-controlled task comparison but not a fixed-budget causal comparison.

A.3 PREREGISTERED 50-SEED COMPONENT ABLATIONS

The table is populated only after all five configurations reach exact Primary-14 $\times$ 50 coverage and the five full-versus-ablation contrasts receive a joint Holm correction. Until then, its pending cells are claim-disabled. In the no-guards arm, a NaN, crash, or out-of-range action is retained as a formal failure rather than silently dropped.

A.4 EARLIER FIVE-SEED DIAGNOSTIC

This screen predates the 50-seed ratio freeze. Its overlapping task-cluster intervals permit only a qualitative observation that CKA and CS can be complementary. It does not substitute for the preregistered Primary-14 $\times$ 50 component ablation containing CKA-only, CS-only, $w_i = 1$, no guards, and no functional adjudication.

Table 5: Preregistered GR00T component ablations on Primary-14 with 50 paired seeds per task. Δ is full GDSQ-VLA minus the ablation; p is Holm-adjusted over five contrasts. Crashes, NaNs, and out-of-range actions in the no-guards arm count as failures.

Config.	SR (%)	Δ [95% CI]	Holm p	Fail. (%)
Full GDSQ-VLA	65.9	—	—	—
CKA only	<i>pending</i>	<i>pending</i>	<i>pending</i>	<i>pending</i>
CS only	<i>pending</i>	<i>pending</i>	<i>pending</i>	<i>pending</i>
Uniform weights ($w_i = 1$)	<i>pending</i>	<i>pending</i>	<i>pending</i>	<i>pending</i>
No guards	<i>pending</i>	<i>pending</i>	<i>pending</i>	<i>pending</i>
No adjudication	<i>pending</i>	<i>pending</i>	<i>pending</i>	<i>pending</i>

Table 6: Five-seed diagnostic ablation. This screen predates the 50-scenario ratio freeze and therefore studies signal composition rather than the final 16:1 mask. Green cells compare quantized configurations only.

Configuration	Held-out 14 SR	All 18 SR
FP16 (uncompressed)	71.4	73.3
W4A8 + ATM/OHB	52.9	52.2
CS only	58.6	61.1
CKA only	57.1	60.0
CKA + ATM/OHB	60.0	63.3
CS+CKA (64:1)	67.1	67.8

A.5 STATIC CORRECTION ABLATION

On the complete GR00T 50-task matrix, static ATM+OHB changes task-macro SR from 50.8% to 49.4%; the paired interval $[-3.28, 0.52]$ points includes zero. On the complete $\pi_{0.5}$ static matrix, ATM+OHB changes 25.3% to 26.6%, while the preregistered held-out interval $[-0.39, 3.13]$ and Holm-adjusted $p = 0.542$ do not establish a gain. Neither combined configuration is an allowed output of the v8 selector. Standalone ATM and OHB remain static ablations; only the selector-chosen OHB runtime path populates the final row, and it does so through a separate complete 2,500-episode manifest.

A.6 EAGER FAKE-QUANT RUNTIME DIAGNOSTIC

The GR00T final plan stores an estimated 1.001 GiB over the LLM+DiT Linear scope, versus 1.993 GiB for FP16, giving $1.99\times$ theoretical compression. The $\pi_{0.5}$ plan stores 1.868 GiB versus 4.113 GiB. The GR00T and $\pi_{0.5}$ Ω -QVLA estimates are 0.599 and 1.307 GiB, respectively, and additionally include required permutation indices and A4 scale tables. These estimates include tightly packed low-bit weights, weight scales, block rotations, and native FP16 for protected layers and biases. They exclude the vision stack, activations, CUDA workspaces, and simulator state. The current GR00T and $\pi_{0.5}$ Ω -QVLA evaluation packs are 6.41 and 34.97 GiB because they store dequantized FP16 weights and dense rotations; they validate the W4A4 numerical path but are not the packed-deployment sizes reported in Table 1.

These measurements diagnose the current implementation rather than support a deployment claim. The eager path retains floating-point weights and executes floating-point GEMMs. On GR00T it takes 0.254 seconds per replan and peaks at 17.7 GiB, compared with 0.139 seconds and 12.8 GiB for FP16. Packed checkpoint size, isolated policy-server memory, and batch-1 latency require fused low-bit kernels.

Table 7: Observed eager fake-quant runtime. These measurements explain why we claim theoretical component storage rather than end-to-end speed or live-memory gains. Green cells compare compressed configurations only; FP16 is an uncompressed reference.

Configuration	Wall (s)	Replan (s)	Peak MiB
FP16 (uncompressed)	26.9	0.139	12,769
QuantVLA W4A8	37.5	0.268	17,973
GDSQ-VLA +ATM/OHB	31.6	0.259	17,719
GDSQ-VLA (Ours)	30.9	0.254	17,671